

Chirp! Project Report

ITU BDSA 2025 Group 19

Hans Voldby Larsen havl@itu.dk

1 Design and Architecture of *Chirp!*

1.1 Domain model

The domain model is a consists of the classes in the core domain of interest, and the relationships between them. The domain model is visually shown in figure 1 as a UML diagram. Only the core entities are shown so other classes of the app concerning eg. infrastructure or web pages are not part of the domain model.

The core domain consists of the entities Author, Cheep and Follow, where an Author can publish multiple Cheeps and each Cheep is authored by exactly one Author.

Authors can follow other authors, which is modeled through the Follow entity representing the follower-followee relationship.

Authentication and authorization are handled via ASP.NET Identity and are therefore not part of the core domain model. The property ApplicationUserId in the class Author (which is used to connect a user identity with an Author) is therefore not included in the Domain model.

1.2 Architecture — In the small

Figure 2 illustrates the internal architecture of the Chirp! application. The Solution follows a layered architecture with a clear separation between **domain**, **application**, **infrastructure** and **presentation**.

Figure 3 complements the onion architecture overview by showing a UML package diagram that maps the architectural layers to the concrete project structure and illustrates the dependency relationships between them. All dependencies between classes and inheritances are not shown to avoid the illustration to become too overcrowded.

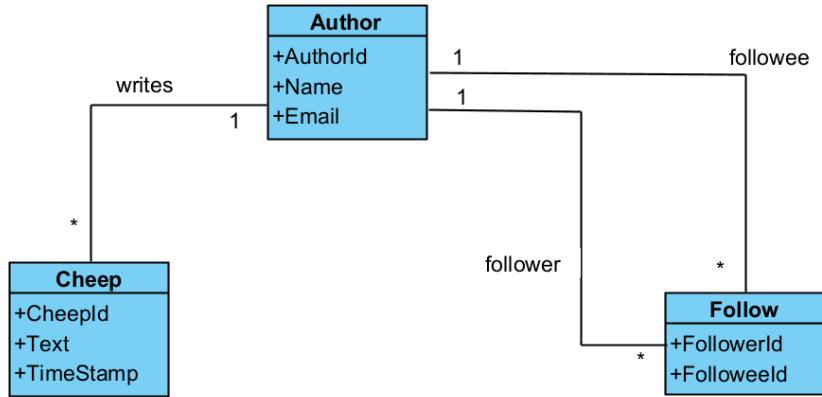


Figure 1: Illustration of the *Chirp!* data model as UML class diagram.

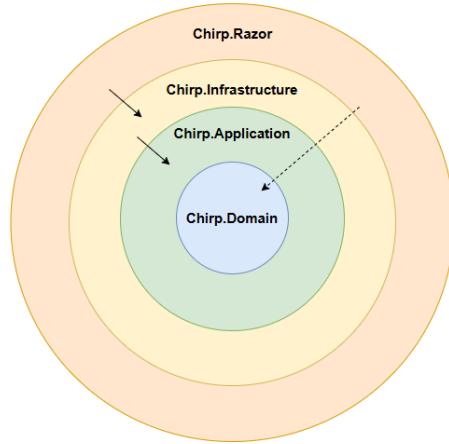


Figure 2: Illustration of the Chirp! onion Archhitecture.

1.2.1 Domain layer

The domain layer (project Chirp.Domain) contains the core domain entities and is completely independent of other layers. The class Author in the domain has a property called ApplicationUserId and is used to associate a domain author with an authenticated user. This property is represented as a primitive value and does not introduce a dependency on the authentication framework.

1.2.2 Application layer

The application layer (project Chirp.Application) defines service interfaces and operates exclusively on data transfer objects (DTOs).

The service interfaces define the operations that the presentation layer can use to interact with the application's functionality. For example, the ICheepService interface specifies methods for retrieving cheeps in different contexts, such as paginated lists of cheeps, cheeps authored by a specific author, cheeps associated with a user, and cheeps forming a user timeline. By letting the presentation layer depend abstract interfaces and not concrete implementations of these services, the presentation layer remains decoupled from the underlying infrastructure, allowing service implementations to be replaced or modified without affecting higher layers.

Data Transfer Objects (DTOs) are used to decouple the application and presentation layers from the domain model by exposing only the data required for a specific use case.

1.2.3 Infrastructure layer

The infrastructure layer (project Chirp.Infrastructure) is responsible for data persistence and contains repositories which encapsulate database access logic and isolate the rest of the application from the Entity Framework Core-specific details. Furthermore it provides the concrete implementations of the service interfaces defined in the application layer.

The data persistence infrastructure is centered around the ChirpDbContext, which serves as the Entity Framework Core database context and defines the mapping between the domain entities and the underlying relational database. The ChirpDbContext is kept lightweight and contains only the aggregate entities from the domain layer (Author, Cheep, and Follow).

The ChirpDbContext inherits from IdentityDbContext, which integrates ASP.NET Identity into the persistence layer and is responsible for managing the database tables related to authentication and authorization (e.g., user accounts, roles, and claims).

Supporting components such as database migrations and the DbInitializer are included to manage schema evolution and initial data seeding.

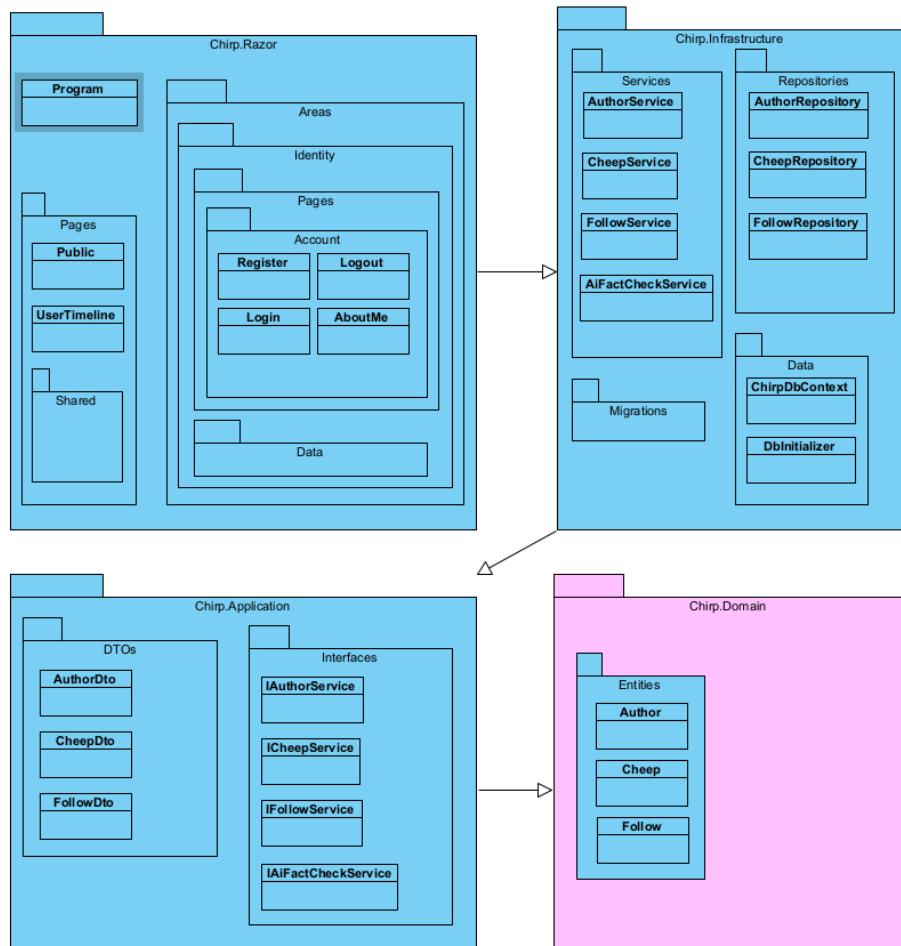


Figure 3: Illustration of the *Chirp!* Archhitecture.

The service implementations (e.g., AuthorService, CheepService, FollowService, and AiFactCheckService) implements the behavior specified by the corresponding application-layer interfaces. These services coordinate domain entities and repositories to fulfill application use cases while remaining hidden behind the abstractions defined in the application layer.

1.2.4 Presentation layer

The presentation layer (project Chirp.Razor) contains all components related to user interaction and request handling. It includes the Razor Pages that render the web-based user interface, the application entry point (Program.cs) where the system is configured and composed, and the authentication-related pages generated through ASP.NET Identity scaffolding. Together, these components are responsible for presenting data to users, handling input, and delegating application logic to the underlying layers.

Ideally, the presentation layer should depend only on service interfaces and data transfer objects (DTOs) defined in the application layer. In practice, some Razor Pages also reference domain entities directly, resulting in limited direct dependencies on the domain layer. While this slightly weakens the intended separation, the overall responsibility boundaries between layers remain clear.

1.2.4.1 Program.cs The Program.cs file defines how the system is assembled at runtime. It is responsible for configuring dependency injection, registering services, setting up authentication and authorization.

Within Program.cs, application-layer service interfaces are bound to their concrete infrastructure-layer implementations, and repositories, database contexts, and identity services are registered with the dependency injection container. The file also configures external authentication via GitHub OAuth and integrates ASP.NET Identity for user management. Also authentication, authorization, static file handling, and session management is configured here. By centralizing these concerns, Program.cs keeps configuration and wiring separate from application logic and presentation concerns.

1.2.4.2 Custom Razor Pages The custom Razor Pages developed for Chirp! implement the core user-facing functionality of the application. These pages handle displaying timelines and creating new cheeps. Each page is responsible for handling HTTP requests, invoking application-layer services, and rendering the resulting data in the user interface.

The Razor Pages primarily interact with the application layer through service interfaces such as ICheepService, IAuthorService, and IFollowService, which encapsulate application logic and data access. In some cases, domain entities are accessed directly from the presentation layer, reflecting pragmatic trade-offs made during development. Nevertheless, the overall flow of control remains from the presentation layer toward the application and infrastructure layers.

1.2.4.3 Razor Pages Generated via ASP.NET Identity Scaffolding

In addition to the custom pages, the presentation layer contains Razor Pages generated using ASP.NET Identity scaffolding. These pages are organized under the Areas/Identity structure and provide functionality related to authentication and account management, including login, logout, registration, and access control.

These scaffolded pages rely on the ASP.NET Identity framework and operate largely independently of the application's domain and application layers. By using the standard scaffolding approach, authentication and authorization concerns are handled using well-established framework components, reducing the need for custom implementation while keeping identity-related functionality isolated from the core application logic.

The “About Me” Page is a custom made Page but is put together with the Identity Pages since it has a tight connection to the user account.

1.3 Architecture of deployed application

Figure X illustrates the deployment architecture of the Chirp! application and shows how the system is composed at runtime. The application is deployed as a client–server web application consisting of a web browser client, a server-side ASP.NET Core application hosted on Azure App Service, a local SQLite database, and external services accessed over HTTPS.

The client side consists of a standard web browser running on a client computer. Users interact with the application through rendered HTML pages and submit requests via forms and links. All communication between the client and the server takes place over HTTPS, ensuring secure transmission of data.

The server side is hosted within an Azure App Service environment and runs a web application that handles incoming HTTP requests, executes application logic, and renders Razor Pages. The server acts as the central coordinator of the system, mediating all communication between the client, the database, and external services.

Data persistence is provided by a SQLite database that is deployed within the same Azure App Service environment as the web application. As SQLite is a file-based database, it is accessed locally by the web application without network communication. The database stores both application-specific data (such as authors, cheeps, and follow relationships) and authentication-related data managed by ASP.NET Identity.

In addition to local persistence, the deployed application integrates with two external services over HTTPS. GitHub OAuth is used to support user authentication via OAuth 2.0, allowing users to sign in using their GitHub accounts. During authentication, the web application communicates directly with GitHub’s OAuth endpoints, while the client browser interacts only with the Chirp! server. Furthermore, the application communicates with the OpenAI API over HTTPS using a REST-based interface to perform AI-assisted fact checking of user-submitted

content. Both external services are accessed exclusively by the server and are not part of the deployed system itself.

Overall, the deployment architecture follows a client–server model in which the browser acts as a client, the application hosted on Azure App Service serves as the central runtime component, and all persistence and external integrations are handled server-side.

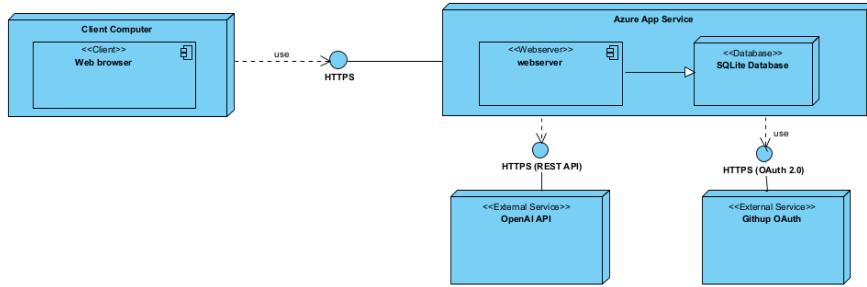


Figure 4: Architecture of the deployd application.

1.4 User activities

1.5 Sequence of functionality/calls trough *Chirp!*

2 Process

2.1 Build, test, release, and deployment

GitHub Actions was used as a means for continuous integration and deployment in the project.

The integration workflow runs on every push and pull request targeting the main branch. It checks the repository, installs .Net 8, restores dependencies, and builds the solution, installs Playwright for the test project, starts the Razor app locally and then runs the test suite. During testing, some configuration is required and is thus provided through environment variables such as OAuth client credentials from Github Secrets, furthermore the application is run in a dedicated testing environment to isolate test behavior from production behavior.

The deployment workflow runs on every push to main and can also be triggered with manually using workflow_dispatch. For deployment, a separate workflow builds and tests the Razor project in Release configuration and then publishes it to a folder which is uploaded as an artifact. A deploy job then downloads said artifact, logs into Azure via azure/login, zip the output, and deploys it to the Azure Web App using azure/webapps-deploy. Separating the build and deploy

steps into two separate jobs ensures that the deployed package is exactly the one that was produced by the build job

For the last workflow, the release workflow, releases are handled with a tag-based workflow triggered by pushing a version / release tag matching the format of 'v..*'. This workflow reruns build and tests as an extra quality gate and then publishes self-contained, single-file builds for Windows, macOS, and Linux. The outputs are zipped and attached to a GitHub Release created automatically.

Together, all of these workflows ensure a coherent structure in which all changes are thoroughly validated before integration, automatically deployed to production, and released for the supported platforms.

All of this can be seen in the UML activity diagrams below

2.2 Team work

2.3 How to make *Chirp!* work locally

2.4 How to run test suite locally

3 Ethics

3.1 License

3.2 LLMs, ChatGPT, CoPilot, and others

LLMs have been used throughout the development of this project for a variety of tasks. The extent to which LLMs were used varied among team members, reflecting different working styles and preferences. Below is an overview of how LLMs contributed to the project.

- **Debugging and troubleshooting.** LLMs (primarily ChatGPT and GitHub Copilot) have been particularly helpful in debugging and troubleshooting. In addition to suggesting fixes, they often explained why errors occurred. This was especially valuable when working with high-level frameworks such as EF Core and ASP.NET Identity, where error messages can be difficult to interpret. Without ChatGPT, resolving these issues would often have required extensive searching across online forums and documentation. Furthermore, GitHub Copilot's built-in analysis of pipeline error logs proved useful when CI builds failed.
- **Writing code.** ChatGPT is very keen on writing code. However, since the primary purpose of this project was educational, prompts were deliberately structured to discourage copy-paste-ready solutions. In some cases, LLM-generated code was still used—for example, when generating unit tests. In such instances, ChatGPT was credited as a co-author in the corresponding commit. In a non-educational or production setting, LLMs would likely be used more extensively for writing code.

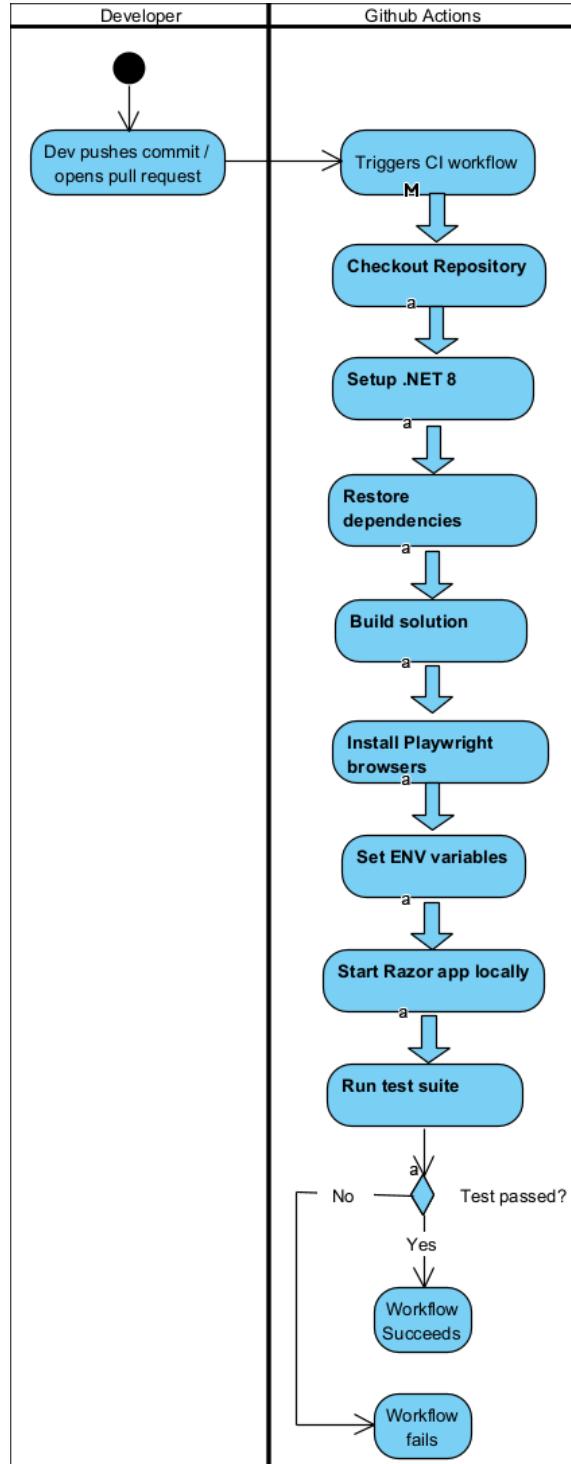


Figure 5: Illustration of the integration workflow as a UML activity diagram.
9

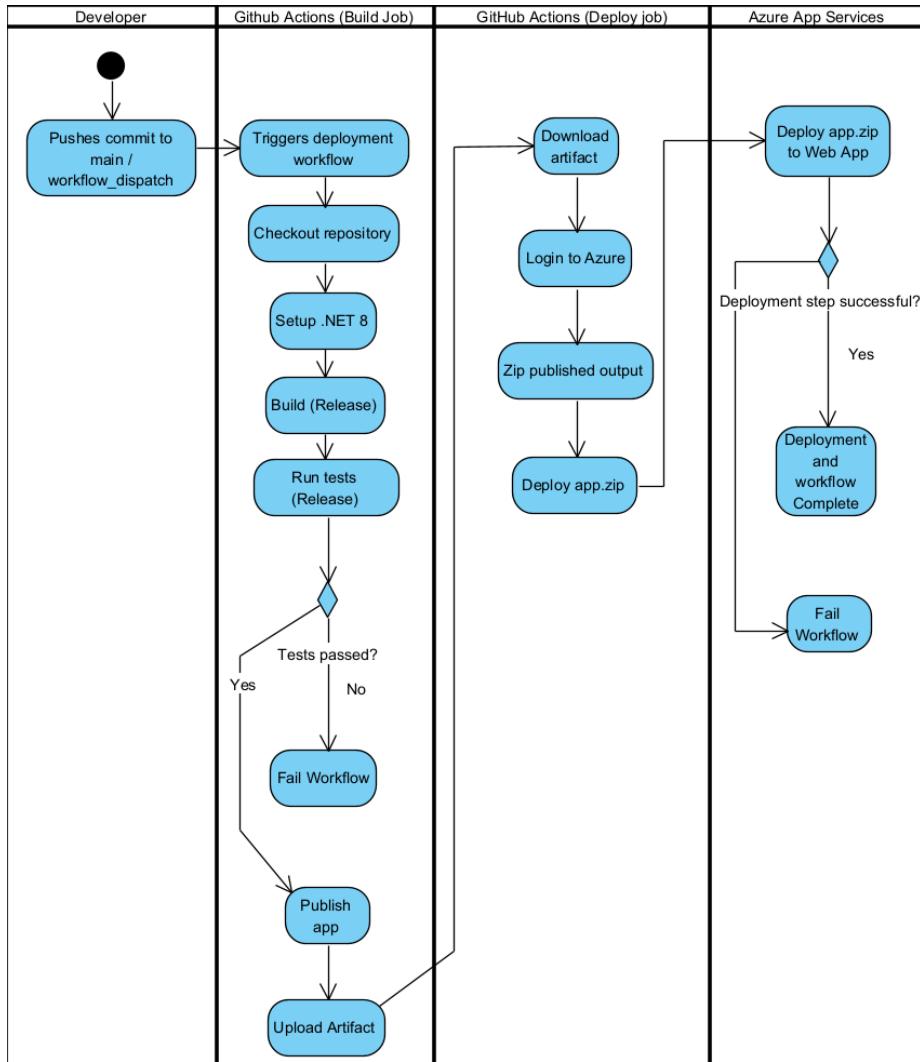


Figure 6: Illustration of the deployment workflow as a UML activity diagram.

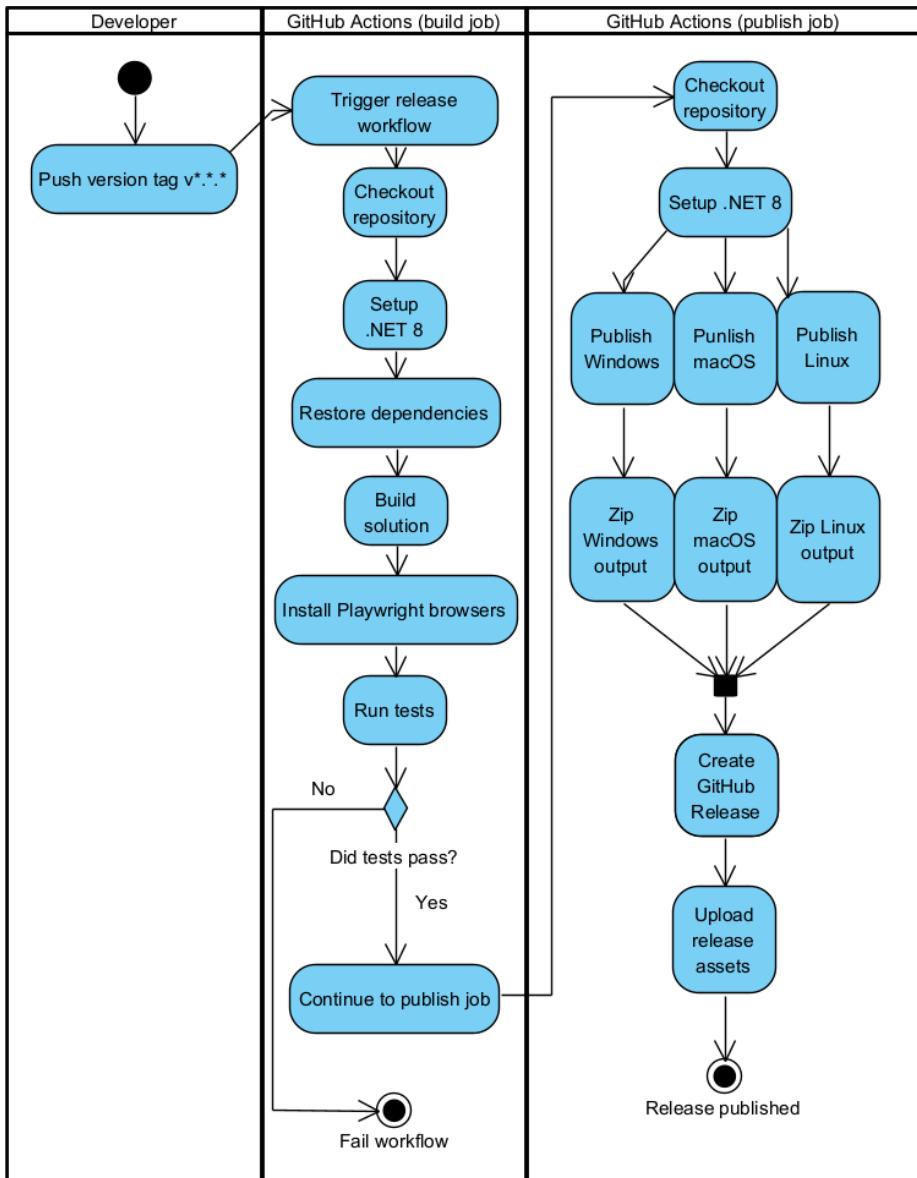


Figure 7: Illustration of the release workflow as a UML activity diagram.

- **Generating commit messages.** In Git Desktop the built-in Copilot has been used to generate commit messages. The commit messages have typically been used as-is or with only a few edits.
- **Personalized tutoring.** ChatGpt has been helpful in explaining concepts, answering follow-up questions, and providing feedback on ideas. The ability to iteratively ask questions and receive explanations tailored to the user's level of understanding significantly accelerated learning.
- **Writing.** Not being native english speakers, ChatGPT has been helpful in reformulating sentences so they appear more professional without "fluff".

There were occasions where ChatGpt provided bad advice. It seems sometimes it fixates on solving a specific problem by creating a very advanced solution and overlooking more obvious ones. But overall the use of LMM has sped up the development and gaining experience with these tools was itself a valuable outcome.

As discussed in the lectures, LLMs consume significant amounts of electricity and water, and their use should therefore be considered within an ethical and environmental context. However, given the growing importance of AI in software development, it is difficult to imagine a future where software developers do not rely on such tools. Choosing not to learn how to use LLMs due to environmental concerns may, in practice, exclude developers from the field altogether. While changes in individual behavior are important, they are unlikely on their own to resolve the environmental impact of LLMs — just as concerns have not eliminated air travel. The more realistic solution will likely come from future technological advances that make AI systems more energy- and resource-efficient.